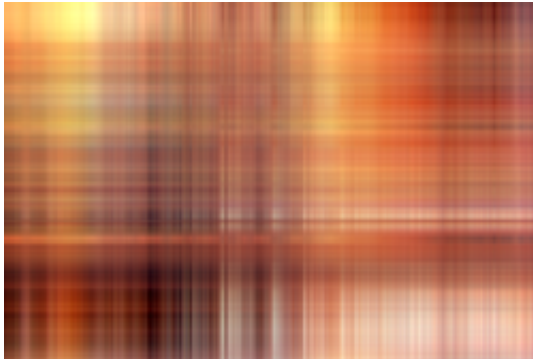


Principal component analysis




Ingo Blechschmidt
December 17th, 2014

Principal component analysis




Ingo Blechschmidt
December 17th, 2014

Principal component analysis



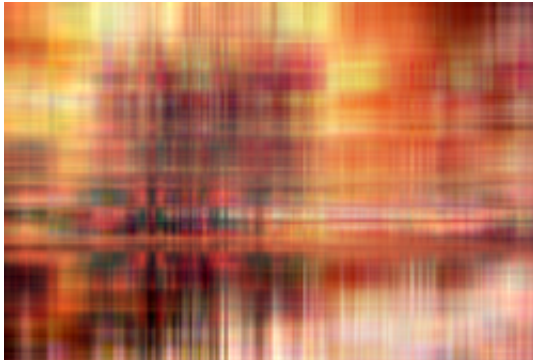
Ingo Blechschmidt
December 17th, 2014

Principal component analysis




Ingo Blechschmidt
December 17th, 2014

Principal component analysis



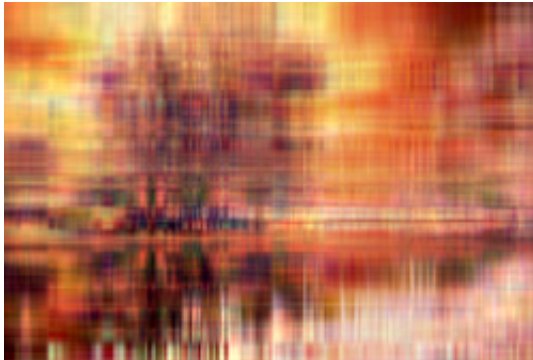
Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



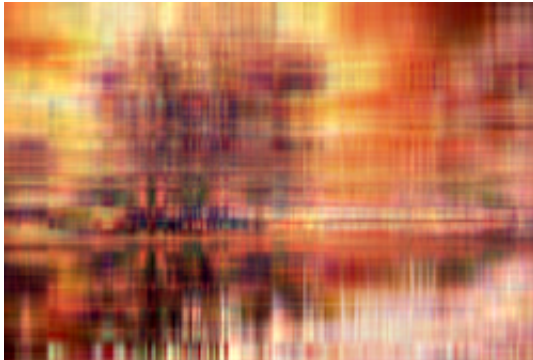
Ingo Blechschmidt
December 17th, 2014

Principal component analysis




Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt


December 17th, 2014

Principal component analysis



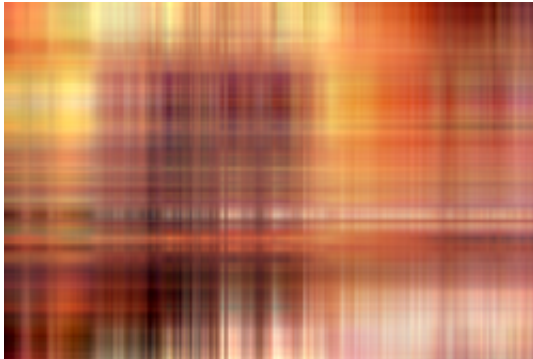
Ingo Blechschmidt
December 17th, 2014

Principal component analysis



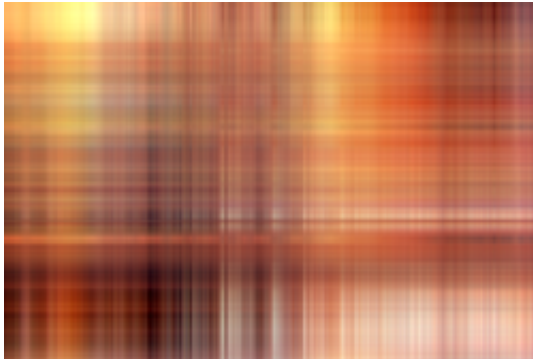
Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

Principal component analysis



Ingo Blechschmidt
December 17th, 2014

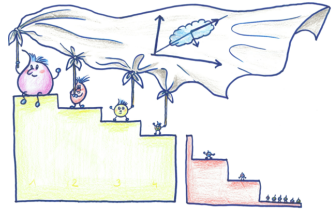
Outline

1 Theory

- Singular value decomposition
- Pseudoinverses
- Low-rank approximation

2 Applications

- Image compression
- Proper orthogonal decomposition
- Principal component analysis
- Eigenfaces
- Digit recognition



Singular value decomposition

Let $A \in \mathbb{R}^{n \times m}$. Then there exist

- numbers $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_m \geq 0$,
- an orthonormal basis $\mathbf{v}_1, \dots, \mathbf{v}_m$ of \mathbb{R}^m , and
- an orthonormal basis $\mathbf{w}_1, \dots, \mathbf{w}_n$ of \mathbb{R}^n ,

such that

$$A\mathbf{v}_i = \sigma_i\mathbf{w}_i, \quad i = 1, \dots, m.$$

In matrix language:

$$A = W\Sigma V^t,$$

where $V = (\mathbf{v}_1 | \dots | \mathbf{v}_m) \in \mathbb{R}^{m \times m}$ orthogonal,

$W = (\mathbf{w}_1 | \dots | \mathbf{w}_n) \in \mathbb{R}^{n \times n}$ orthogonal,

$$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_m) \in \mathbb{R}^{n \times m}.$$

- The singular value decomposition (SVD) exists for any real matrix, even rectangular ones.
- The singular values σ_i are unique.
- The basis vectors are not unique.
- If A is orthogonally diagonalizable with eigenvalues λ_i (for instance, if A is symmetric), then $\sigma_i = |\lambda_i|$.
- $\|A\|_{\text{Frobenius}} = \sqrt{\sum_{ij} A_{ij}^2} = \sqrt{\text{tr}(A^t A)} = \sqrt{\sum_i \sigma_i^2}$.
- There exists a generalization to complex matrices. In this case, the matrix A can be decomposed as $W\Sigma V^*$, where V^* is the complex conjugate of V^t and W and V are unitary matrices.
- The singular value decomposition can also be formulated in a basis-free manner as a result about linear maps between finite-dimensional Hilbert spaces.

Existence proof (sketch):

1. Consider the eigenvalue decomposition of the symmetric and positive-semidefinite matrix $A^t A$: We have an orthonormal basis \mathbf{v}_i of eigenvectors corresponding to eigenvalues λ_i .
2. Set $\sigma_i := \sqrt{\lambda_i}$.
3. Set $\mathbf{w}_i := \frac{1}{\sigma_i} A \mathbf{v}_i$ (for those i with $\lambda_i \neq 0$).
4. Then $A \mathbf{v}_i = \sigma_i \mathbf{w}_i$ holds trivially.
5. The \mathbf{w}_i are orthonormal: $(\mathbf{w}_i, \mathbf{w}_j) = \frac{1}{\sigma_i \sigma_j} (A^t A \mathbf{v}_i, \mathbf{v}_j) = \frac{\lambda_i \delta_{ij}}{\sigma_i \sigma_j}$.
6. If necessary, extend the \mathbf{w}_i to an orthonormal basis.

This proof gives rise to an algorithm for calculating the SVD, but unless $A^t A$ is small, it has undesirable numerical properties. (But note that one can also use AA^t !) Since the 1960ies, there exists a stable iterative algorithm by Golub and van Loan.

The pseudoinverse of a matrix

Let $A \in \mathbb{R}^{n \times m}$ and $\mathbf{b} \in \mathbb{R}^n$. Then the solutions to the optimization problem

$$\|A\mathbf{x} - \mathbf{b}\|_2 \longrightarrow \min$$

under $\mathbf{x} \in \mathbb{R}^m$ are given by

$$\mathbf{x} = A^+ \mathbf{b} + V \begin{pmatrix} 0 \\ \star \end{pmatrix},$$

where $A = W\Sigma V^t$ is the SVD and

$$\begin{aligned} A^+ &= W\Sigma^+ V^t, \\ \Sigma^+ &= \text{diag}(\sigma_1^{-1}, \dots, \sigma_m^{-1}). \end{aligned}$$

- In the formula for Σ^+ , set $0^{-1} := 0$.
- If A happens to be invertible, then $A^+ = A^{-1}$.
- The pseudoinverse can be used for polynomial approximation: Let data points $(x_i, y_i) \in \mathbb{R}^2$, $1 \leq i \leq N$, be given. Want to find a polynomial $p(z) = \sum_{k=0}^n \alpha_k z^k$, $n \ll N$, such that

$$\sum_{i=1}^N |p(x_i) - y_i|^2 \longrightarrow \min.$$

In matrix language, this problem is written

$$\|A\mathbf{u} - \mathbf{y}\|_2 \longrightarrow \min$$

where $\mathbf{u} = (\alpha_0, \dots, \alpha_N)^T \in \mathbb{R}^{n+1}$ and

$$A = \begin{pmatrix} 1 & x_1 & x_1^2 & \cdots & x_1^n \\ 1 & x_2 & x_2^2 & \cdots & x_2^n \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_N & x_N^2 & \cdots & x_N^n \end{pmatrix} \in \mathbb{R}^{N \times (n+1)}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix} \in \mathbb{R}^N.$$

Low-rank approximation

Let $A = W\Sigma V^t \in \mathbb{R}^{n \times m}$ and $1 \leq r \leq n, m$. Then a solution to the optimization problem

$$\|A - M\|_{\text{Frobenius}} \longrightarrow \min$$

under all matrices M with $\text{rank } M \leq r$ is given by

$$M = W\Sigma_r V^t,$$

$$\text{where } \Sigma_r = \text{diag}(\sigma_1, \dots, \sigma_r, 0, \dots, 0).$$

The approximation error is

$$\|A - W\Sigma_r V^t\|_F = \sqrt{\sigma_{r+1}^2 + \dots + \sigma_m^2}.$$

- This is the Eckart–Young(–Mirsky) theorem.
- Beware of false and incomplete proofs in the literature!

Image compression

- Think of images as matrices.
- Substitute a matrix $W\Sigma V^t$ by $W\Sigma_r V^t$ with r small.
- To reconstruct $W\Sigma_r V^t$, only need to know
 - the r singular values $\sigma_1, \dots, \sigma_r$, r
 - the first r columns of W , and height $\cdot r$
 - the top r rows of V^t . width $\cdot r$
- Total amount:
 $r \cdot (1 + \text{height} + \text{width}) \ll \text{height} \cdot \text{width}$

- See <http://speicherleck.de/iblech/stuff/pca-images.pdf> for sample compressions and <http://pizzaseminar.speicherleck.de/skript4/08-principal-component-analysis-svd-image.py> for the Python code producing these images.
- Image compression by singular value decomposition is mostly of academic interest only.
- This might be for the following reasons: other compression algorithms have more efficient implementations; other algorithms tailor to the specific properties of human vision; the basis vectors of other approaches (for instance, DCT) are similar to the most important singular basis vectors of a sufficiently large corpus of images.
- See <http://dsp.stackexchange.com/questions/7859/relationship-between-dct-and-pca>.

Proper orthogonal decomposition

Given data points $\mathbf{x}_i \in \mathbb{R}^N$, want to find a low-dimensional linear subspace which **approximately contains** the \mathbf{x}_i .

Minimize

$$J(U) := \sum_i \|\mathbf{x}_i - P_U(\mathbf{x}_i)\|^2$$

under all r -dimensional subspaces $U \subseteq \mathbb{R}^N$, $r \ll N$,
where $P_U : \mathbb{R}^N \rightarrow \mathbb{R}^N$ is the orthogonal projection onto U .

Proper orthogonal decomposition

Given data points $\mathbf{x}_i \in \mathbb{R}^N$, want to find a low-dimensional linear subspace which **approximately contains** the \mathbf{x}_i .

Minimize

$$J(U) := \sum_i \|\mathbf{x}_i - P_U(\mathbf{x}_i)\|^2$$

under all r -dimensional subspaces $U \subseteq \mathbb{R}^N$, $r \ll N$,
where $P_U : \mathbb{R}^N \rightarrow \mathbb{R}^N$ is the orthogonal projection onto U .

More concrete formulation: Minimize

$$J(\mathbf{u}_1, \dots, \mathbf{u}_r) := \sum_i \left\| \mathbf{x}_i - \sum_{j=1}^r \langle \mathbf{x}_i, \mathbf{u}_j \rangle \mathbf{u}_j \right\|^2,$$

where $\mathbf{u}_1, \dots, \mathbf{u}_r \in \mathbb{R}^N$, $\langle \mathbf{u}_j, \mathbf{u}_k \rangle = \delta_{jk}$.

- In the first formulation, the optimization domain is the *Grassmannian* of r -dimensional subspaces in \mathbb{R}^N . It is a compact topological space (in fact a manifold of dimension $r \cdot (N - r)$). Since $J(U)$ depends continuously on U , the optimization problem is guaranteed to have a solution.
- The solution is in general not unique, not even locally. For instance, consider the four data points $(\pm 1, \pm 1)$ in \mathbb{R}^2 . Then any line U through the origin solves the optimization problem, with functional value $J(U) = 4$.
- In the more concrete formulation, we look for an orthonormal basis of a suitable subspace. In this case, the optimization domain is a compact subset of $\mathbb{R}^{N \times r}$.
- Since a given subspace possesses infinitely many orthonormal bases (at least for $r \geq 2$), solutions to this refined problem are never unique, not even locally.
- Note that this is a non-convex optimization problem. Therefore common numerical techniques do not apply.

Collect the data points \mathbf{x}_i as columns of a matrix

$$X = (\mathbf{x}_1 | \cdots | \mathbf{x}_\ell) \in \mathbb{R}^{N \times \ell}$$

and consider its singular value decomposition

$$X = W\Sigma V^t.$$

Then a solution to the minimization problem is given by the first r columns of W , with approximation error

$$J = \sum_i \|\mathbf{x}_i\|^2 - \sum_{j=1}^r \sigma_j^2.$$

- Proper orthogonal decomposition (POD) cannot be used to find low-dimensional *submanifolds* which approximately contain given data points. But check out *kernel principal component analysis*.
- Also, POD does not work well with *affine* subspaces. But in this case, the fix is easy: Simply shift the data points so that their mean is zero.
- POD is a general method for *dimension reduction* and can be used as a kind of “preconditioner” for many other algorithms: Simply substitute the given points \mathbf{x}_i by their projections $P_U(\mathbf{x}_i)$.

Principal component analysis

Given observations $x_i^{(k)}$ of random variables $X^{(k)}$, want to find **linearly uncorrelated** principal components.

Write $X = (\mathbf{x}_1 | \cdots | \mathbf{x}_\ell) \in \mathbb{R}^{N \times \ell}$. Calculate $X = W \Sigma V^t$.
Then the principal components are the variables

$$Y^{(j)} = \sum_k W_{kj} X^{(k)}.$$

Most of the variance is captured by $Y^{(1)}$; second to most is captured by $Y^{(2)}$; and so on.


- For instance, in a study about circles, the variables *radius*, *diameter*, and *circumference* are linearly correlated.
- Principal component analysis (PCA) would automatically pick one of these attributes as a principal component.
- We have to normalize the data to have zero empirical mean first. Then XX^t is the empirical covariance matrix.
- Note that, in the given sample, the $Y^{(j)}$ are indeed uncorrelated:

$$\begin{aligned}
 E(Y^{(j)}XX^tY^{(k)}) &= (W\mathbf{e}_j)^tXX^t(W\mathbf{e}_k) \\
 &= \mathbf{e}_j^tW^tW\Sigma V^tV\Sigma^tW^tW\mathbf{e}_k \\
 &= \Sigma\Sigma^t.
 \end{aligned}$$

- Beware that PCA cannot resolve nonlinear correlation.
- Also note that PCA is sensitive to outliers and not scaling-independent.


Eigenfaces

- Record sample faces
 $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^{\text{width} \cdot \text{height}}$.
- Calculate a POD basis of **eigenfaces**.
- Recognize faces by looking at the coefficients of the most important eigenfaces.



Eigenfaces resemble faces.

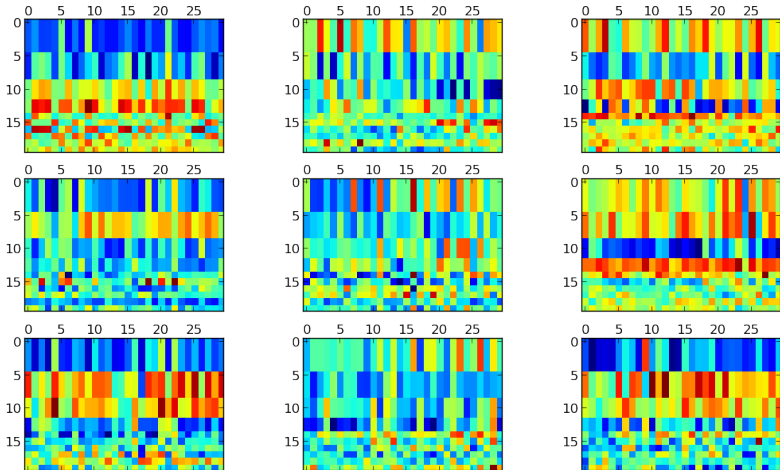
More eigenfaces



- A naive approach is very sensitive to lighting, scale and translation.
- But extensions are possible, for instance considering the eyes, the nose, and the mouth separately; this leads to eigeneyes, eigennoses, and eigenmouths.
- Image credit:
<http://upload.wikimedia.org/wikipedia/commons/6/67/Eigenfaces.png>
<http://www.cenparmi.concordia.ca/~jdong/eigenface.gif>
- Live demo:
<http://cognitron.psych.indiana.edu/nsfgrant/FaceMachine/faceMachine.html>
- Examples:
<http://www.cs.princeton.edu/~cdecoro/eigenfaces/>

Digit recognition

Apply POD for dimension reduction, then use some similarity measure or clustering technique. Results:




- These images were produced by a Python program. The actual numerical code is very short (a few lines). Of course, the MNIST data set was used.

<http://pizzaseminar.speicherleck.de/skript4/08-principal-component-analysis/digit-recognition.py>

- The nine images show the values of the first ten POD coefficients of the first 30 samples of the digits 1 to 9. Each column corresponds to a different sample. The first four POD coefficients are drawn using more vertical space, so that visual weight aligns with importance.
- One can clearly see that the POD coefficients differ for the different digits.
- Also, one can see that the difference is not so great for similar digits like 5 and 8 or 7 and 9.

Eigendigits



- These images show the first 12 POD basis vectors. The first basis vector is a kind of “prototypical digit”. The other basis vectors give subsequent “higher-order terms”.
- Because I didn’t implement a similarity measure or clustering technique, I couldn’t calculate the percentage of correctly classified digits. However, presumably the success rate would not be too high: Like the eigenfaces approach, this naive implementation is sensitive to the specific position of the digits in the bounding box. Refined techniques are discussed in the literature.

**I don't always do
model reduction**

**but when I do I use
singular value decomposition**